





Offset direction and size change the role-score pattern

My saved activations · layer 12 · five-role probe · one gardening conversation

User direction



Tool direction



—●— User

—●— CoT

—●— Assistant

Each point averages the same source-role tokens; 512 matched tokens in total. Lines join five evaluated offsets.

Offline arithmetic only: I edit saved states and rescore them. No downstream model pass or behavior is measured.