



Probe quality varies by role, depth and split

My H100 run · five-role probe · all 24 layers · 249 base texts in the full corpus

Prompt split: 124 held-out variants



Text split: 24 held-out base texts



— System — User — CoT — Assistant — Tool

Token-weighted held-out recall. These splits use different evaluation sets; their difference is descriptive.