

±1% offsets along saved User and Tool probe directions

GPT-OSS-20B · layer 12 · five-role probe · one conversation · 512 matched tokens · all User tags



Passage key: U = User, C = CoT (reasoning), A = Assistant; 1 and 2 identify the two turns of each role.

Reference norm = 45.2471, the median across 921 forwarded tokens; 1% / 5% give lengths 0.4525 / 2.2624.

Every token and both directions use the same fixed offset length at each strength. Colors retain the paper's role meanings.

Offline rescoring; no downstream model computation or behavior. "Subtract" applies an opposite offset, not component removal.