





A small Tool offset raises both scores for original User text

0 → +1% Tool: Userness +3.2 points; Toolness +3.3 points (before rounding the means)

User direction



Tool direction



● User ▲ CoT ■ Assistant

GPT-OSS-20B · layer 12 · five-role probe · one conversation; each mean uses the same source-role tokens.

Reference norm = 45.2471 across 921 forwarded tokens. Lines join five evaluated strengths; no intermediate values measured.

Offline rescoring of 512 saved states; no model continuation or behavior. Marker shapes also identify source roles.