


My saved gardening activations: $\pm 1\%$ offsets along User and Tool probe directions

Offline probe-response illustration · gpt-oss-20b · layer 12 · five-role probe · all user tags



Colors identify the original passages: User, CoT and Assistant, following the paper. All panels retain the same 512 tokens.

I edit saved activations and score them with the same probe. These plots do not show downstream model or behavior changes.

“Subtract” means an opposite offset, not erasing a role. Both directions use the same offset norm; strengths are illustrative.