


"Toolward" steering was pointing at chain of thought

Class-mean role directions from the probe corpus (9 base texts $\times$ 5 roles, dev-10) · filled points at the probe site, open points at block 11



A direction is the difference between mean activations of the same passages placed in two role wrappers, e.g. Tool minus User.

The Tool-User vector used in earlier steering work is almost the CoT-User vector. Adding it raises the CoT reading about as much as the Tool reading.

The trained probe rows give the same picture (directions/geometry.json). Cosines are properties of the corpus geometry, not of any behavior.